

Hadoop & HBase & Phoenix install and connect

- 서버 6대 가이드를 계속 진행한다는 가정
- Hadoop 3.1.2
- centos 7
- zookeeper 3.4.10
- HBase 2.2.5

Oracle VM VirtualBox 관리자

파일(F) 머신(M) 도움말(H)

도구

새로 만들기(N) 설정(S) 삭제 표시(H)

이름	상태
an01 (hbase-phoenix-spark)	실행 중
sn01 (hbase-phoenix-spark)	실행 중
rm01 (hbase-phoenix-spark)	실행 중
dn01 (hbase-phoenix-spark)	실행 중
dn02 (hbase-phoenix-spark)	실행 중
dn03 (hbase-phoenix-spark)	실행 중

일반

이름: sn01
운영 체제: Red Hat (64-bit)

시스템

기본 메모리: 4096 MB
프로세서: 2
부팅 순서: 플로피, 광 디스크, 하드 디스크
가속: VT-X/AMD-V, 네스티드 페이징, PAE/NX, KVM 반가상화

미리 보기

일반

이름: rm01
운영 체제: Red Hat (64-bit)

시스템

기본 메모리: 4096 MB
프로세서: 2
부팅 순서: 플로피, 광 디스크, 하드 디스크
가속: VT-X/AMD-V, 네스티드 페이징, PAE/NX, KVM 반가상화

미리 보기

1. HBase install

- HBase 설치
- Hadoop 계정에서 운영
- an01 서버에서 설정 후 데이터 노드들에게 배포
- an01: master, dn01, dn02, dn03: region server

action server: an01
pwd: /home/hadoop

wget <http://archive.apache.org/dist/hbase/2.2.5/hbase-2.2.5-bin.tar.gz>

mv hbase-2.2.5 hbase // 폴더명 변경

cd hbase/conf 경로 변경

1. HBase install

- HBase 설정

```
action server: an01  
pwd: /home/hadoop/hbase/conf  
  
vi hbase-site.xml
```

```
<?xml version="1.0"?>  
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>  
<configuration>  
  <property>  
    <name>hbase.rootdir</name>  
    <value>hdfs://NNHA/user/hbase</value>  
  </property>  
  
  <property>  
    <name>hbase.master</name>  
    <value>an01:6000</value>  
  </property>  
  
  <property>  
    <name>hbase.cluster.distributed</name>  
    <value>true</value>  
  </property>
```

```
<property>
  <name>hbase.zookeeper.quorum</name>
  <value>an01,sn01,rm01</value>
</property>

<property>
  <name>dfs.replication</name>
  <value>1</value>
</property>

<property>
  <name>hbase.zookeeper.property.clientPort</name>
  <value>2181</value>
</property>

<property>
  <name>hbase.zookeeper.property.dataDir</name>
  <value>/home/hadoop/data/zookeeper</value>
</property>

<property>
  <name>dfs.datanode.max.xcievers</name>
  <value>4096</value>
</property>
<property>
  <name>hbase.unsafe.stream.capability.enforce</name>
  <value>false</value>
</property>
</configuration>
```

1. HBase install

- 하둡 설정 파일 복사

```
action server: an01
pwd: /home/hadoop/hbase/conf

cp /home/hadoop/hadoop-3.1.2/etc/hadoop/hdfs-site.xml /home/hadoop/hbase/conf/hdfs-site.xml
cp /home/hadoop/hadoop-3.1.2/etc/hadoop/core-site.xml /home/hadoop/hbase/conf/core-site.xml
```

- **region servers** 설정

```
action server: an01
pwd: /home/hadoop/hbase/conf

vi regionservers

dn01
dn02
dn03
```

1. HBase install

- hbase 설정

action server: an01

pwd: /home/hadoop/hbase/conf

vi hbase-env.sh

java 경로 설정

regionservers 파일 경로 설정

HBASE_MANAGES_ZK = false 설정 // 직접 구축한 주키퍼 환경에서 실행하겠다는 의미

```
# The java implementation to use.  Java 1.8+ required.  
export JAVA_HOME=/opt/apps/jdk8
```

```
# File naming hosts on which HRegionServers will run.  $HBASE_HOME/conf/regionservers by default.  
export HBASE_REGIONSERVERS=/home/hadoop/hbase/conf/regionservers
```

```
# Tell HBase whether it should manage it's own instance of ZooKeeper or not.  
export HBASE_MANAGES_ZK=false
```

1. HBase install

- 하둡 jar 파일 HBase/lib 복사

```
action server: an01
pwd: /home/hadoop/hbase/conf

vi hbase-env.sh

cp /home/hadoop/hadoop-3.1.2/share/hadoop/yarn/timelineservice/lib/htrace-core-3.1.0-incubating.jar
/home/hadoop/hbase/lib/htrace-core-3.1.0-incubating.jar

cd ~ // 경로 변경
```

- 재압축 후 배포

```
action server: an01
pwd: /home/hadoop

tar xvfz hbase.tar.gz hbase-2.2.5

scp hbase.tar.gz hadoop@dn01:/home/hadoop
scp hbase.tar.gz hadoop@dn02:/home/hadoop
scp hbase.tar.gz hadoop@dn03:/home/hadoop
```


1. HBase install

- 압축 해제

```
action server: dn01, dn02, dn03
```

```
pwd: /home/hadoop
```

```
tar xvfz hbase.tar.gz
```

```
exit // root 계정 전환
```

- HBase 환경 변수 설정 및 폴더 생성

```
action server: an01, dn01, dn02, dn03
```

```
pwd: /root
```

```
vi /etc/profile.d/hbase.sh
```

```
export HBASE_HOME=/home/hadoop/hbase
```

```
export PATH=$PATH:$HBASE_HOME/bin
```

```
source /etc/profile.d/hbase.sh
```

```
su - hadoop // 하둡 계정 재접속
```

```
mkdir data
```

```
cd data
```

```
mkdir zookeeper
```

```
cd /home/hadoop/hbase
```

1. HBase install

- HBase 실행
- 정상 실행 시 HMaster, HRegionServer가 실행됨

action server: an01
pwd: /home/hadoop/hbase

./bin/start-hbase.sh
jps // 상태 확인

an01

```
[hadoop@an01 hbase]$ jps
1520 JournalNode
1682 DFSZKFailoverController
8807 HMaster
9048 Jps
1598 NameNode
[hadoop@an01 hbase]$
```

dn01

```
[hadoop@dn01 hbase]$ jps
1921 DataNode
3761 Jps
1429 NodeManager
3445 HRegionServer
[hadoop@dn01 hbase]$
```

dn02

```
[hadoop@dn02 hbase]$ jps
1952 DataNode
1426 NodeManager
3893 HRegionServer
4219 Jps
[hadoop@dn02 hbase]$
```

dn03

```
[hadoop@dn03 hbase]$ jps
1425 NodeManager
3713 HRegionServer
4025 Jps
1917 DataNode
[hadoop@dn03 hbase]$
```

- **웹 확인**

http://192.168.56.100:16010

Master: an01

←

→

↺

⚠ 주의 요함

192.168.56.100:16010/master-status

🔍

☆

⚙

한글

⋮

APACHE

HBASE

Home

Table Details

Procedures & Locks

HBCk Report

Process Metrics

Local Logs

Log Level

Debug Dump

Metrics Dump

Profiler

HBase Configuration

Master an01

Region Servers

Base Stats

Memory

Requests

Storefiles

Compactions

Replications

ServerName	Start time	Last contact	Version	Requests Per Second	Num. Regions
dn01,16020,1623732969088	Tue Jun 15 13:56:09 KST 2021	0 s	2.2.5	0	13
dn02,16020,1623732968925	Tue Jun 15 13:56:08 KST 2021	0 s	2.2.5	0	16
dn03,16020,1623732968567	Tue Jun 15 13:56:08 KST 2021	0 s	2.2.5	0	12
Total:3				0	41

1. HBase install

- csv 파일 insert 테스트
- 지하철 예제 데이터가 an01:/home/hadoop/testdata 로컬에 저장되어있음을 가정
- csv 파일을 삽입할 빈 테이블 생성

```
action server: an01
```

```
pwd: /home/hadoop/hbase
```

```
hdfs dfs -mkdir /user/batch_data // hdfs 폴더 생성
```

```
hdfs dfs -put /home/hadoop/testdata/subway2015.csv /user/batch_data // 로컬 csv -> hdfs
```

```
hbase shell // hbase 명령어 모드
```

```
create "SUBWAY2015", {NAME => "cf2015"} // 컬럼 패밀리가 cf2015인 테이블 생성
```

```
list // 테이블 확인
```

```
exit // 명령어 모드 종료
```

```
hbase(main):001:0> list
TABLE
SUBWAY2015
```

1. HBase install

- csv 파일 insert

action server: an01

pwd: /home/hadoop/hbase

```
./bin/hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -Dimporttsv.separator="," -  
Dimporttsv.columns="HBASE_ROW_KEY,cf2015:use_dt,cf2015:line_num,cf2015:sub_sta_nm,cf2015:ride_pasgr_num,cf2015:alight_pasgr_num,cf2015:work_dt" SUBWAY2015 /user/batch_data/subway2015.csv
```

1번째 줄: ImportTsv 임포트, 쉼표로 구분됨을 정의

2, 3번째 줄: 컬럼들을 정의, HBase 특성 상 고유한 row key가 있어야 함. 예제 csv파일에서는 이미 인덱스 열이 존재. 앞서 hbase shell을 통해 테이블 지정시 cf2015라는 컬럼 패밀리를 지정하였고, 컬럼 패밀리 내부의 컬럼들의 정의 후 테이블 명과 삽입할 hdfs 상의 csv파일의 경로를 정의

후 맵 리듀스 잡과 함께 데이터가 삽입

hbase shell // 명령어 모드 접속

scan 'SUBWAY2015', {LIMIT => 5} // SUBWAY2015 테이블의 5행만 출력, HBase의 특성 상 바이트 형식으로 출력된다.

```
hbase(main):002:0> scan 'SUBWAY2015', {LIMIT => 5}
COLUMN+CELL
1 column=cf2015:alight_pasgr_num, timestamp=1623734540155, value=40197
1 column=cf2015:line_num, timestamp=1623734540155, value=1\xED\x98\xB8\xEC\x84\xA0
1 column=cf2015:ride_pasgr_num, timestamp=1623734540155, value=47071
1 column=cf2015:sub_sta_nm, timestamp=1623734540155, value=\xEC\x84\x9C\xEC\x9A\xB8\xEC\x97\xAD
1 column=cf2015:use_dt, timestamp=1623734540155, value=20150101
1 column=cf2015:work_dt, timestamp=1623734540155, value=20151217
10 column=cf2015:alight_pasgr_num, timestamp=1623734540155, value=5654
10 column=cf2015:line_num, timestamp=1623734540155, value=1\xED\x98\xB8\xEC\x84\xA0
10 column=cf2015:ride_pasgr_num, timestamp=1623734540155, value=5345
10 column=cf2015:sub_sta_nm, timestamp=1623734540155, value=\xEB\x8F\x99\xEB\xAC\x98\xEC\x95\x9E
10 column=cf2015:use_dt, timestamp=1623734540155, value=20150101
10 column=cf2015:work_dt, timestamp=1623734540155, value=20151217
100 column=cf2015:alight_pasgr_num, timestamp=1623734540155, value=7836
100 column=cf2015:line_num, timestamp=1623734540155, value=4\xED\x98\xB8\xEC\x84\xA0
100 column=cf2015:ride_pasgr_num, timestamp=1623734540155, value=8432
100 column=cf2015:sub_sta_nm, timestamp=1623734540155, value=\xEB\xAF\xB8\xEC\x95\x84
100 column=cf2015:use_dt, timestamp=1623734540155, value=20150101
100 column=cf2015:work_dt, timestamp=1623734540155, value=20151217
1000 column=cf2015:alight_pasgr_num, timestamp=1623734540155, value=7414
1000 column=cf2015:line_num, timestamp=1623734540155, value=\xEB\xB6\x84\xEB\x8B\xB9\xEC\x84\xA0
1000 column=cf2015:ride_pasgr_num, timestamp=1623734540155, value=6443
1000 column=cf2015:sub_sta_nm, timestamp=1623734540155, value=\xEA\xB0\x80\xEC\xB2\x9C\xEB\x8C\x80
1000 column=cf2015:use_dt, timestamp=1623734540155, value=20150102
1000 column=cf2015:work_dt, timestamp=1623734540155, value=20151217
10000 column=cf2015:alight_pasgr_num, timestamp=1623734540155, value=55118
10000 column=cf2015:line_num, timestamp=1623734540155, value=7\xED\x98\xB8\xEC\x84\xA0
10000 column=cf2015:ride_pasgr_num, timestamp=1623734540155, value=53647
10000 column=cf2015:sub_sta_nm, timestamp=1623734540155, value=\xEA\xB0\x80\xEC\x82\xB0\xEB\x94\x94\xEC\xA7\x80\xED\x84\xB8\xEB\x8B\xA8\xEC\xA7\x80
10000 column=cf2015:use_dt, timestamp=1623734540155, value=20150119
10000 column=cf2015:work_dt, timestamp=1623734540155, value=20151217
5 row(s)
Took 0.0454 seconds
hbase(main):003:0> █
```

1. HBase install

- 바이트 형식의 데이터 출력

```
scan "SUBWAY2015",{COLUMNS=>'cf2015:line_num:toString',LIMIT=>5} // 저장된 데이터 형식에 맞게 튜닝
```

```
hbase(main):007:0> scan "SUBWAY2015",{COLUMNS=>'cf2015:line_num:toString',LIMIT=>5}
ROW                                COLUMN+CELL
 1                                column=cf2015:line_num, timestamp=1623734540155, value=1호 선
10                                column=cf2015:line_num, timestamp=1623734540155, value=1호 선
100                               column=cf2015:line_num, timestamp=1623734540155, value=4호 선
1000                             column=cf2015:line_num, timestamp=1623734540155, value=분 당 선
10000                           column=cf2015:line_num, timestamp=1623734540155, value=7호 선

5 row(s)
Took 0.0085 seconds
hbase(main):008:0>
```

1. HBase install

- 테이블 삭제
- HBase의 테이블을 삭제하기 위해서는 먼저 테이블을 비활성화 해야함
- 다음 Phoenix와의 연동 테스트를 위해 삭제 하지 않음.

```
disable "SUBWAY2015" // 테이블 비활성화  
drop "SUBWAY2015" // 테이블 삭제
```

```
list // 삭제 확인
```

```
exit // 명령어 모드 종료
```

- HBase 종료

```
action server: an01  
pwd: /home/hadoop/hbase
```

```
./bin/stop-hbase.sh
```


2. Phoenix install

- HBase를 SQL로 활용하기 위한 phoenix 연동
- Hbase 서버를 종료한 채 진행
- HMaster 서버에 설치

```
action server: an01  
pwd: /home/hadoop
```

```
wget https://mirror.navercorp.com/apache/phoenix/apache-phoenix-5.0.0-HBase-2.0/bin/apache-phoenix-5.0.0-HBase-2.0-bin.tar.gz
```

```
tar xvfz apache-phoenix-5.0.0-HBase-2.0-bin.tar.gz
```

```
mv /apache-phoenix-5.0.0-HBase-2.0-bin phoenix // 폴더명 변경
```

```
cd phoenix // 경로 변경
```

2. Phoenix install

- **phoenix-5.0.0-HBase-2.0-client.jar**
- **phoenix-5.0.0-HBase-2.0-server.jar**
- **phoenix-core-5.0.0-HBase-2.0.jar**
- **모든 HBase 클러스터의 lib 폴더에 복사**

```
action server: an01
```

```
pwd: /home/hadoop/phoenix
```

```
cp ./phoenix-5.0.0-HBase-2.0-client.jar /home/hadoop/hbase/lib/phoenix-5.0.0-HBase-2.0-client.jar
```

```
cp ./phoenix-5.0.0-HBase-2.0-server.jar /home/hadoop/hbase/lib/phoenix-5.0.0-HBase-2.0-server.jar
```

```
cp ./phoenix-core-5.0.0-HBase-2.0.jar /home/hadoop/hbase/lib/phoenix-core-5.0.0-HBase-2.0.jar
```

```
scp phoenix-5.0.0-HBase-2.0-server.jar hadoop@dn01:/home/hadoop/hbase/lib
```

```
scp phoenix-5.0.0-HBase-2.0-client.jar hadoop@dn01:/home/hadoop/hbase/lib
```

```
scp phoenix-core-5.0.0-HBase-2.0.jar hadoop@dn01:/home/hadoop/hbase/lib
```

```
scp phoenix-5.0.0-HBase-2.0-server.jar hadoop@dn02:/home/hadoop/hbase/lib
```

```
scp phoenix-5.0.0-HBase-2.0-client.jar hadoop@dn02:/home/hadoop/hbase/lib
```

```
scp phoenix-core-5.0.0-HBase-2.0.jar hadoop@dn03:/home/hadoop/hbase/lib
```

```
scp phoenix-5.0.0-HBase-2.0-server.jar hadoop@dn03:/home/hadoop/hbase/lib
```

```
scp phoenix-5.0.0-HBase-2.0-client.jar hadoop@dn03:/home/hadoop/hbase/lib
```

```
scp phoenix-core-5.0.0-HBase-2.0.jar hadoop@dn03:/home/hadoop/hbase/lib
```

2. Phoenix install

- phoenix 환경 변수 설정

```
action server: an01
```

```
pwd: /root
```

```
vi /etc/profile.d/phoenix.sh
```

```
export PHOENIX_HOME=/home/hadoop/phoenix
```

```
export PATH=$PATH:$HADOOP_HOME/bin
```

```
export PATH=$PATH:$HADOOP_HOME/sbin
```

```
source /etc/profile.d/phoenix.sh
```

```
su - hadoop
```

```
cd phoenix
```

2. Phoenix install

- 하둡, HBase 설정 파일을 phoenix/bin에 복사

```
action server: an01
```

```
pwd: /home/hadoop/phoenix
```

```
cp /home/hadoop/hadoop-3.1.2/etc/hadoop/core-site.xml /home/hadoop/bin/phoenix/core-site.xml
```

```
cp /home/hadoop/hadoop-3.1.2/etc/hadoop/hdfs-site.xml /home/hadoop/bin/phoenix/hdfs-site.xml
```

```
cp /home/hadoop/hbase/conf/hbase-site.xml /home/hadoop/phoenix/bin/hbase-site.xml
```

```
cd .. // 경로 변경
```

- hbase, phoenix 실행

```
action server: an01
```

```
pwd: /home/hadoop
```

```
cd hbase
```

```
./bin/start-hbase.sh
```

```
cd ~
```

```
cd phoenix
```

```
./bin/sqlline.py // 명령어 모드 접속
```

- print error 시: python 버전 다운, phoenix는 python2와 호환됨

```
!tables // 테이블 확인
```

- phoenix 실행 시 system과 관련된 테이블이 자동 생성됨

```
0: jdbc:phoenix:> !tables
```

TABLE_CAT	TABLE_SCHEM	TABLE_NAME	TABLE_TYPE	REMARKS	TYPE_NAME	SELF_REFERENCING_COL_NAME	REF_GENERATION	INDEX_STATE	IMMUTABLE_ROWS	SALT_BUCKETS	MULTI_TENANT
	SYSTEM	CATALOG	SYSTEM TABLE						false	null	false
	SYSTEM	FUNCTION	SYSTEM TABLE						false	null	false
	SYSTEM	LOG	SYSTEM TABLE						true	32	false
	SYSTEM	SEQUENCE	SYSTEM TABLE						false	null	false
	SYSTEM	STATS	SYSTEM TABLE						false	null	false

```
0: jdbc:phoenix:> 
```

- hbase web을 통해 자동 생성된 phoenix system 테이블이 hbase 테이블에도 생성되면 연동되었음을 확인(<http://192.168.56.100:16010>)

Master: an01

192.168.56.100:16010/master-status

HomeTable DetailsProcedures & LocksHBase ReportProcess MetricsLocal LogsLog LevelDebug DumpMetrics DumpProfilerHBase Configuration

Tables

User TablesSystem TablesSnapshots

6 table(s) in set. [Details]

Namespace	Name	State		Regions									Description
				OPEN	OPENING	CLOSED	CLOSING	OFFLINE	FAILED	SPLIT	Other		
default	SYSTEM.CATALOG	ENABLED	1	0	0	0	0	0	0	0	0	'SYSTEM CATALOG', (TABLE_ATTRIBUTES => (PRIORITY => '2000', coprocessor\$1 => 'org.apache.phoenix.coprocessor.ScanRegionObserver(805306366)', coprocessor\$2 => 'org.apache.phoenix.coprocessor.UngroupedAggregateRegionObserver(805306366)', coprocessor\$3 => 'org.apache.phoenix.coprocessor.GroupedAggregateRegionObserver(805306366)', coprocessor\$4 => 'org.apache.phoenix.coprocessor.ServerCachingEndpointImpl(805306366)', coprocessor\$5 => 'org.apache.hadoop.hbase.coprocessor.MultiRowMutationEndpoint(805306366)', coprocessor\$6 => 'org.apache.phoenix.coprocessor.MetaDataEndpointImpl(805306366)', coprocessor\$7 => 'org.apache.phoenix.coprocessor.MetaDataRegionObserver(805306367)', METADATA => ('SPLIT_POLICY' => 'org.apache.phoenix.schema.MetaDataSplitPolicy')), (NAME => '0', DATA_BLOCK_ENCODING => 'FAST_DIFF', BLOOMFILTER => 'NONE')	
default	SYSTEM.FUNCTION	ENABLED	1	0	0	0	0	0	0	0	0	'SYSTEM FUNCTION', (TABLE_ATTRIBUTES => (PRIORITY => '2000', coprocessor\$1 => 'org.apache.phoenix.coprocessor.ScanRegionObserver(805306366)', coprocessor\$2 => 'org.apache.phoenix.coprocessor.UngroupedAggregateRegionObserver(805306366)', coprocessor\$3 => 'org.apache.phoenix.coprocessor.GroupedAggregateRegionObserver(805306366)', coprocessor\$4 => 'org.apache.phoenix.coprocessor.ServerCachingEndpointImpl(805306366)', coprocessor\$5 => 'org.apache.phoenix.hbase.index.Indexer(805306366)Index.Builder=org.apache.phoenix.index.PhoenixIndexBuilder.org.apache.hadoop.hbase.index.codec.class=org.apache.phoenix.index.PhoenixIndexCodec', coprocessor\$6 => 'org.apache.phoenix.coprocessor.MetaDataEndpointImpl(805306367)', METADATA => ('SPLIT_POLICY' => 'org.apache.phoenix.schema.SystemFunctionSplitPolicy')), (NAME => '0', DATA_BLOCK_ENCODING => 'FAST_DIFF', BLOOMFILTER => 'NONE')	
default	SYSTEM.LOG	ENABLED	32	0	0	0	0	0	0	0	0	'SYSTEM LOG', (TABLE_ATTRIBUTES => (PRIORITY => '2000', coprocessor\$1 => 'org.apache.phoenix.coprocessor.ScanRegionObserver(805306366)', coprocessor\$2 => 'org.apache.phoenix.coprocessor.UngroupedAggregateRegionObserver(805306366)', coprocessor\$3 => 'org.apache.phoenix.coprocessor.GroupedAggregateRegionObserver(805306366)', coprocessor\$4 => 'org.apache.phoenix.coprocessor.ServerCachingEndpointImpl(805306366)', coprocessor\$5 => 'org.apache.phoenix.hbase.index.Indexer(805306366)Index.Builder=org.apache.phoenix.index.PhoenixIndexBuilder.org.apache.hadoop.hbase.index.codec.class=org.apache.phoenix.index.PhoenixIndexCodec', coprocessor\$6 => 'org.apache.phoenix.coprocessor.MetaDataEndpointImpl(805306367)', METADATA => ('SPLIT_POLICY' => 'org.apache.phoenix.schema.SystemFunctionSplitPolicy')), (NAME => '0', DATA_BLOCK_ENCODING => 'FAST_DIFF', TTL => '604800 SECONDS (7 DAYS)', BLOOMFILTER => 'NONE')	
default	SYSTEM.MUTEX	ENABLED	1	0	0	0	0	0	0	0	0	'SYSTEM MUTEX', (NAME => '0', TTL => '900 SECONDS (15 MINUTES)')	
default	SYSTEM.SEQUENCE	ENABLED	1	0	0	0	0	0	0	0	0	'SYSTEM SEQUENCE', (TABLE_ATTRIBUTES => (PRIORITY => '2000', coprocessor\$1 => 'org.apache.phoenix.coprocessor.ScanRegionObserver(805306366)', coprocessor\$2 => 'org.apache.phoenix.coprocessor.UngroupedAggregateRegionObserver(805306366)', coprocessor\$3 => 'org.apache.phoenix.coprocessor.GroupedAggregateRegionObserver(805306366)', coprocessor\$4 => 'org.apache.phoenix.coprocessor.ServerCachingEndpointImpl(805306366)', coprocessor\$5 => 'org.apache.phoenix.hbase.index.Indexer(805306366)Index.Builder=org.apache.phoenix.index.PhoenixIndexBuilder.org.apache.hadoop.hbase.index.codec.class=org.apache.phoenix.index.PhoenixIndexCodec', coprocessor\$6 => 'org.apache.phoenix.coprocessor.SequenceRegionObserver(805306366)', (NAME => '0', DATA_BLOCK_ENCODING => 'FAST_DIFF', BLOOMFILTER => 'NONE')	
default	SYSTEM.STATS	ENABLED	1	0	0	0	0	0	0	0	0	'SYSTEM STATS', (TABLE_ATTRIBUTES => (PRIORITY => '2000', coprocessor\$1 => 'org.apache.phoenix.coprocessor.ScanRegionObserver(805306366)', coprocessor\$2 => 'org.apache.phoenix.coprocessor.UngroupedAggregateRegionObserver(805306366)', coprocessor\$3 => 'org.apache.phoenix.coprocessor.GroupedAggregateRegionObserver(805306366)', coprocessor\$4 => 'org.apache.phoenix.coprocessor.ServerCachingEndpointImpl(805306366)', coprocessor\$5 => 'org.apache.hadoop.hbase.coprocessor.MultiRowMutationEndpoint(805306366)', METADATA => ('SPLIT_POLICY' => 'org.apache.phoenix.schema.SystemStatsSplitPolicy')), (NAME => '0', DATA_BLOCK_ENCODING => 'FAST_DIFF', BLOOMFILTER => 'NONE')	

2. Phoenix install

- hbase table을 view로 가져오기

sqlline.py 모드에서

```
CREATE VIEW " SUBWAY2015 " (  
  ROWKEY VARCHAR PRIMARY KEY,  
  " cf2015 " . " use_dt " VARCHAR,  
  " cf2015 " . " line_num " VARCHAR,  
  " cf2015 " . " sub_sta_nm " VARCHAR,  
  " cf2015 " . " ride_pasgr_num " VARCHAR,  
  " cf2015 " . " alight_pasgr_num " VARCHAR,  
  " cf2015 " . " work_dt " VARCHAR);
```

```
select * from SUBWAY2015 limit 5;  
select count(*) from SUBWAY2015;
```

- HBase와 Phoenix의 연동으로 SQL을 사용할 수 있음

```
0: jdbc:phoenix:> select * from SUBWAY2015 limit 5;
```

ROWKEY	use_dt	line_num	sub_sta_nm	ride_pasgr_num	alight_pasgr_num	work_dt
1	20150101	1호 선	서 울 역	47071	40197	20151217
10	20150101	1호 선	동 모 앞	5345	5654	20151217
100	20150101	4호 선	미 아	8432	7836	20151217
1000	20150102	분 당 선	가 천 대	6443	7414	20151217
10000	20150119	7호 선	가 산 디 지 털 단 지	53647	55118	20151217

```
5 rows selected (0.2 seconds)
0: jdbc:phoenix:> select count(*) from SUBWAY2015;
```

COUNT(1)
199380

```
1 row selected (1.49 seconds)
0: jdbc:phoenix:>
0: jdbc:phoenix:>
```